Beyond Stochastic Intelligence

Deterministic Governance and the Structural Limits of Scale in AI Systems

Armonti Du-Bose-Hill
Independent Researcher

Paper Type: Systems governance position paper with regulatory translation appendices.

Abstract

Modern artificial intelligence systems have achieved unprecedented capability through scale, yet remain structurally unsuitable as decision authorities in high-stakes environments. This limitation is not a transient artifact of model quality or training data. It is a fundamental architectural constraint. Probabilistic generative systems cannot, by design, provide deterministic guarantees, auditable causality, or enforceable authority.

The prevailing industry response emphasizes additional scale, layered controls, and post hoc filtering. These approaches treat observed failures as implementation defects rather than as structural properties of probabilistic generation. The result is a widening gap between capability and governability as these systems are increasingly positioned closer to execution authority.

The dominant paradigm is constrained by a category error. Generation is repeatedly conflated with governance. Drawing from safety-critical systems engineering, formal methods, and bounded rationality theory, this paper proposes a constraint-first model in which stochastic systems generate hypotheses while deterministic layers retain decision authority.

The analysis explains why scale alone cannot resolve these limitations, examines the failure modes of after-the-fact governance, and presents design principles for governance-native intelligence systems capable of auditability, replayability, and regulatory alignment without reliance on model internals or proprietary mechanisms.

Decision authority in high-stakes systems requires three non-negotiable properties: determinism, replayability, and externally verifiable causality. Any system

that cannot reproduce identical governance outcomes for identical inputs, reconstruct the admissibility path after execution, and expose causal justification independent of narrative explanation cannot serve as a decision authority under consequence. This paper treats these properties as architectural invariants rather than aspirational goals.

1 The Current Landscape: Scale Without Control

Enterprise and public deployments increasingly integrate probabilistic systems into workflows with direct legal, financial, and physical consequence. Capability has outpaced control, and scale has not closed the gap. Training objectives optimize plausibility under uncertainty rather than correctness under constraint. As a result, increasingly capable models expand the failure surface while leaving decision authority undefined.

This conclusion reflects consistent patterns observed across more than one hundred structured interviews conducted with professionals spanning healthcare, finance, logistics, mobility, public sector operations, legal services, and enterprise technology. Across domains, confidence in model capability grows faster than confidence in accountability. Organizations repeatedly report that the core problem is not intelligence, but governability once outputs influence real decisions.

Three properties recur in production systems independent of parameter count, dataset size, or interface design.

Non-determinism as a first-class property

Repeated inference under partially specified context reveals that identical inputs do not guarantee identical outputs when stochastic selection remains part of the inference path. This behavior is not an implementation flaw. It is the mechanism by which these systems trade off diversity, exploration, and uncertainty resolution.

Once outputs carry consequence, this property becomes operationally problematic. Audit, incident response, and root-cause analysis degrade when the system cannot be replayed precisely. Accountability depends on reproducible causality, not approximate reconstruction.

Unverifiable reasoning chains

Intermediate reasoning artifacts are not proof objects and cannot be treated as stable evidence. Narrative explanations cannot serve as verification because they are generated as text rather than derived as constrained derivations. The generator is optimized to produce sequences that resemble reasoning, not to produce verifiable steps that satisfy external invariants.

Governance cannot be built on introspection. Introspection is not a reliable substrate for correctness, particularly when incentives favor fluency over constraint satisfaction.

Authority collapse: generation treated as decision

Outputs are increasingly wired into automation, approvals, and operational pathways that produce real effects. Many deployments implicitly treat generated content as a decision, even when the generator was never designed to hold authority. A recommendation becomes operationally binding once it is consumed by humans under time pressure or fed into downstream automation.

The most damaging failures are not incorrect statements in isolation. They are incorrect statements that cross an execution boundary without deterministic evaluation.

Scale amplifies this pattern. Increased capability does not establish authority. It increases the cost of failure.

2 Why Hallucinations Are a Structural Outcome

Probabilistic generation under incomplete information produces hallucinations as a structural outcome. Systems optimized for plausibility over partially observed distributions behave this way because interpolation and extrapolation are intrinsic to likelihood-optimized generation, and fluency is rewarded even when ground truth is absent.

Treating hallucinations as an implementation defect directs effort toward tuning rather than architecture. The core issue is not that the system occasionally produces unsupported content. The issue is that the system lacks enforceable mechanisms to prevent unsupported content from influencing action.

The Generator–Judge Separation

One component produces candidate hypotheses while another component decides what is admissible to act upon. Generation is probabilistic and exploratory. Judgment is deterministic and constrained. Action under consequence requires enforceable invariants, and enforceable invariants require deterministic evaluation at an authority boundary.

Any architecture that allows the generator to self-certify will eventually produce unbounded failure modes under distribution shift, adversarial inputs, and operational drift.

This separation is not merely conceptual. The inevitability of hallucination follows directly from how probabilistic generators assign likelihood when constraints are incomplete.

Let x denote the prompt or context, y a generated output, and K a reference knowledge base representing the authoritative grounding substrate declared for the task and decision context. Define a support predicate $S(y, x, K) \in \{0, 1\}$ where $S = 1$ if and only if every factual claim in y is supported by K or explicitly contained in x.

The probability of hallucination under a generative policy $p_\theta(y \mid x)$ can then be written as:

$$P_\theta(\text{hallucination} \mid x) = \mathbb{E}_{y \sim p_\theta(y|x)} \left[1 - S(y, x, K)\right]$$

If the model assigns nonzero probability mass to unsupported outputs, the probability of hallucination is strictly greater than zero. Because probabilistic generators assign likelihood prior to admissibility evaluation, this probability cannot be driven to zero through scale, decoding strategy, or training objective alone. Hallucination is therefore not an anomaly but an irreducible structural outcome whenever generation precedes deterministic constraint.

This formulation makes hallucination measurable without reliance on model internals. Unsupported outputs become observable governance events rather than semantic failures, and without an external, deterministic admissibility check, such outputs remain structurally permitted to propagate.

Consequence Is the Signal

Failures surface only when outputs carry consequence. The primary risk is not incorrect generation, but ungoverned acceptance. Retrospective review degrades under load, and governance that depends on sustained human attention fails at scale.

Error containment belongs at the execution boundary. Remediation after execution is not governance. It is forensics.

3 The Hidden Cost of Reactive Governance

Enterprise deployment patterns that wrap generators with filters, policy checks, and human review gates share a structural weakness. Post hoc governance is reactive, and reactive governance cannot guarantee prevention. Evaluation occurs after an output has been produced and after it has entered a decision pathway. Reactive governance creates the appearance of control while leaving authority structurally undefined.

Reactive Authority Is Ineffective Authority

A violation is detected only after a candidate action has been proposed or communicated. Detection after generation is not prevention. The violation already exists in the operational environment. Controls that intervene after emission are strictly weaker than controls that govern before execution. Mature systems move governance upstream into deterministic admission control.

Policies as Advisory Rather Than Enforceable

When policy logic is implemented as a soft layer around a generator rather than as a hard gate in the execution path, it can be bypassed through integration gaps, exception paths, and predictable user behavior under pressure. The generator is not bound to policy as an invariant.

As systems become operationally critical, organizational incentives increasingly favor throughput over enforcement. Governance becomes a cost center that is optimized away unless it is architecturally unavoidable.

Audit Trails That Log Outputs Rather Than Causality

Logging that captures only the final output does not record the chain of admissibility decisions that allowed it to cross into execution. Output logging without causal logging is insufficient for regulated domains that require traceability.

Traceability requires reconstructing why an action was permitted. Auditability depends on deterministic governance layers that emit explicit, versioned decision artifacts rather than narrative outputs.

4 Intelligence as a Constrained Dynamical System

Operational behavior is a property of the full stack. System design, not model design, determines how intelligence under consequence behaves. Intelligence under consequence emerges from constrained orchestration rather than unconstrained generation. Operational behavior is therefore governed by constraint placement , not by the model's expressiveness.

Reliability, auditability, and bounded behavior require enforceable invariants evaluated outside the generator to avoid self-certification. Once systems cross into execution authority, architecture becomes the primary lever.

Constraints as First-Class Citizens

Constraints are executable rules, not preferences. Instructions are interpreted by the generator. Constraints are evaluated by a separate authority layer with explicit semantics. They must reside in code paths that are independently testable, versionable, and reviewable.

Deterministic Governance as an External Boundary

A probabilistic subsystem proposes actions that may be useful, unsafe, or invalid. Deterministic governance is the only stable mechanism for enforcing invariants at an execution boundary.

Deterministic mechanisms can be tested, replayed, and audited. Probabilistic mechanisms cannot provide equivalent guarantees under identical conditions.

Bounded Rationality as the Theoretical Basis

Decision-making under finite compute, incomplete data, and irreducible uncertainty reveals that bounded rationality is a systems constraint that shapes architecture. Constrained orchestration emerges when exhaustive evaluation is impossible.

Deterministic Governance Beyond the Execution Boundary

Systems that allow probabilistic outputs to cross execution boundaries without deterministic admission control cannot guarantee auditability, traceability, or bounded failure under scale. Deterministic governance is therefore a necessary system property once probabilistic systems influence execution.

5 Determinism Does Not Mean Rigidity

Hybrid system design allows exploration to remain probabilistic while execution remains constrained. Determinism belongs to governance, not generation. Systems can remain flexible internally while bounded externally.

6 Separation of Concerns: Generator vs. Governor

Authority design requires explicit separation. Generators produce candidates. Governors enforce admissibility. Entanglement creates self-justification and collapses auditability. Separation is the structural precondition for inspection and control.

7 A Deprecated Direction: Training-Time Constraint Embedding

Embedding governance inside the generator reintroduces opacity. Training-time behavior does not guarantee inference-time compliance. Differentiable approximations of policy are not enforceable rules. Governance must remain external, legible, and versioned.

8 Regulatory Tailwinds and Architectural Pressure

Regulatory frameworks increasingly emphasize traceability, documentation, and lifecycle governance. These properties align with deterministic constraints and explicit control structures. Architectures that treat governance as an afterthought encounter escalating friction.

9 Enterprise Deployment: Patterns That Survive Contact With Reality

Governance scales only when centralized as a platform capability. Distributed control logic produces drift and bypass paths. This reflects the structural approach already used for identity, access control, and logging in mature infrastructures, extended to probabilistic components.

10 Civilizational-Scale Implications

Systems that influence infrastructure, finance, healthcare, mobility, and law cannot tolerate non-determinism at execution boundaries. Remediation requires causal reconstruction, not narrative explanation.

11 Conclusion: Intelligence Is a System Property

Applied intelligence is not a property of models in isolation, but of systems that combine generation, constraints, memory, and governance into a coherent execution path. Once probabilistic components influence real-world outcomes, architectural choices determine whether intelligence remains assistive or becomes authoritative.

Probabilistic generators can propose, summarize, and explore. They cannot serve as decision authorities under consequence. Decision authority requires deterministic evaluation, replayable outcomes, and externally verifiable causality. Systems that collapse this boundary trade short-term capability gains for long-term governability loss.

This shift is neither optional nor theoretical. Organizations will either design for governed intelligence deliberately, or converge on it later through regulatory pressure, incident response, and accumulated failure. The constraint is architectural, and the choice is temporal.

References

Clements, P. et al. (2010). *Documenting Software Architectures: Views and Beyond*. Addison-Wesley.

European Union. (2024). Regulation (EU) 2024/1689 (Artificial Intelligence Act).

ISO/IEC. (2023). *ISO/IEC 42001:2023 Artificial intelligence – Management system*.

Ji, S. et al. (2023). Survey on hallucination in large language models. *arXiv:2311.05232*.

Leveson, N. (2011). *Engineering a Safer World*. MIT Press.

NIST. (2023). *AI Risk Management Framework 1.0*.

Simon, H. (1955). A behavioral model of rational choice. *Quarterly Journal of Economics*.

Woodcock, J. et al. (2009). Formal methods: Practice and experience. *ACM Computing Surveys*.

Zhang, Y. et al. (2023). Siren's song in the AI ocean. *arXiv:2309.01219*.